

EDUCATION

Adani University <i>B.Tech Computer Science Engineering – CGPA: 8.50/10</i>	Ahmedabad, India <i>Expected Jun 2026</i>
---	--

CERTIFICATIONS

Introduction to Model Context Protocol & Advanced Topics <i>issued by Anthropic</i>	Dec 2025
AWS Academy Graduate - AWS Academy Cloud Developing <i>issued by Amazon Web Services Training and Certification</i>	Nov 2024

TECHNICAL SKILLS

Languages: Python, TypeScript, JavaScript, SQL
Agentic AI: LLM Integration (OpenSource LLMs), A2A, RAG Pipelines, Embeddings, MCP, Multi-Agent Orchestration
Frameworks: React, Next.js, FastAPI, Express.js, Hono, Electron.js, Tailwind CSS, Zod
Data: PostgreSQL, pgvector, Redis, BM25, Vector Search
Cloud & DevOps: AWS, Docker, CI/CD, Cloudflare, NixOS, Git

WORK EXPERIENCE

ProductSquads A2A, LLMs, Agentic RAG, MCP, Agentic AI <i>Associate AI Engineer</i>	Jan 2025 – Present
- Won Jury Special Award at Praveg × ProductSquads Codeathon for an A2A-compatible Intelligent Document Processing Agent. - Productionized the IDP agent on AWS Lambda using GPT-4V for automated OCR, document classification, and metadata extraction. - Built a containerized financial data-extraction pipeline (Python, FastAPI) with a taxonomy-mapping engine that reconstructs complex tables into semantic HTML, achieving high recall across 200+ accounting concepts at 90%+ field-level accuracy. - Implemented server-side persistence across distributed streaming Lambda agents powering the AI Assistant micro-frontend and its BFF, resolving stateful conversation management in an event-driven architecture.	
Meru Fintech Electron.js, Microservices, gRPC, TypeScript, React <i>Full Stack Intern</i>	Jun 2025 – July 2025
- Engineered a full-stack Tally integration solution with a desktop application and backend microservice, solving Tally’s lack of a cloud-based API for hellobooks.ai. - Architected a custom OAuth2-like authentication framework with PKCE to ensure secure communication between the web, desktop client, and backend services. - Established CI/CD pipeline with automated builds and releases; built developer CLI tools for configuration management and real-time diagnostics, optimizing deployment workflows.	

PROJECTS

Deep Research Agent <i>Autonomous Multi-Agent System, Recursive Research Loops, Chain-of-Thoughts</i>	github.com/Hetav21/deep-research-agent
- Engineered an autonomous multi-agent system using Python and FastAPI that performs recursive web research and synthesis. - Implemented an asynchronous research pipeline with configurable depth/breadth controls, utilizing DuckDuckGo for search and BeautifulSoup for high-performance content scraping. - Implemented a scalable orchestration layer using OpenAI Agents SDK to coordinate specialized Query, Search, and Synthesis agents for complex problem-solving tasks.	
GDPR Compliance Intelligence Platform <i>Hybrid Retrieval (Dense+Sparse), RRF, Multi-Query Expansion, Multi-Modal ETL</i>	github.com/Hetav21/gdpr-rag
- Architected a Multi-Modal RAG system combining dense vector search (pgvector) + BM25 sparse retrieval with Reciprocal Rank Fusion, improving legal query relevance over single-retrieval baselines. - Built a high-fidelity ETL pipeline using unstructured and GPT-4o Vision to ingest complex documents, automatically generating searchable metadata for non-text elements like charts and tables. - Implemented an advanced Multi-Query Expansion layer and scalable async architecture using Redis and pgvector, significantly improving recall and system throughput.	